

## Objective:

The main objective of this project is to design and implement a robust data preprocessing system that addresses common challenges such as missing values, outliers, inconsistent formatting, and noise. By performing effective data preprocessing, the project aims to enhance the quality, reliability, and usefulness of the data for machine learning.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.feature_selection import VarianceThreshold

from sklearn.preprocessing import StandardScaler, MinMaxScaler

import warnings
warnings.filterwarnings('ignore') # to prevent warning msgs
```

## IMPORTING DATASET

```
dataset=pd.read_csv('Employee.csv')
```

```
df=pd.DataFrame(dataset)
print("Original dataset:")
print(df.head())
```

Original dataset:

	Company	Age	Salary	Place	Country	Gender
0	TCS	20.0	NaN	Chennai	India	0
1	Infosys	30.0	NaN	Mumbai	India	0
2	TCS	35.0	2300.0	Calcutta	India	0
3	Infosys	40.0	3000.0	Delhi	India	0
4	TCS	23.0	4000.0	Mumbai	India	0

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 148 entries, 0 to 147
```

```
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	Company	140 non-null	object
1	Age	130 non-null	float64
2	Salary	124 non-null	float64
3	Place	134 non-null	object
4	Country	148 non-null	object
5	Gender	148 non-null	int64

```
dtypes: float64(2), int64(1), object(3)
memory usage: 7.1+ KB
```

```
df.describe()
```

	Age	Salary	Gender
count	130.000000	124.000000	148.000000
mean	30.484615	5312.467742	0.222973
std	11.096640	2573.764683	0.417654
min	0.000000	1089.000000	0.000000
25%	22.000000	3030.000000	0.000000
50%	32.500000	5000.000000	0.000000
75%	37.750000	8000.000000	0.000000
max	54.000000	9876.000000	1.000000

## Data Cleaning: (Score : 2)

Find the missing and inappropriate values, treat them appropriately.

Remove all duplicate rows.

Find the outliers.

Replace the value 0 in age as NaN

Treat the null values in all columns using any measures (removing/ replace the values with mean/median/mode)

## Finding Missing Data and Duplicates

```
df.isnull()
```

	Company	Age	Salary	Place	Country	Gender
0	False	False	True	False	False	False
1	False	False	True	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False
..	...	...	...	...	...	...
143	False	False	False	False	False	False
144	False	False	False	False	False	False
145	False	False	False	False	False	False
146	False	False	False	False	False	False

```

147      False  False  False  False  False  False
[148 rows x 6 columns]
df.isnull().sum()
Company      8
Age          18
Salary       24
Place        14
Country       0
Gender        0
dtype: int64

df['Company'].isnull().sum()
8

df['Company'].fillna('Unknown', inplace=True)
df['Company'].isnull().sum()
0

```

## Replace the value 0 in age as NaN

```

print("Number of zeros in 'Age' before replacement:", (df['Age'] ==
0).sum())
df['Age'] = df['Age'].replace(0, np.NaN)
print("Number of NaN values in 'Age' after replacement:",
df['Age'].isna().sum()) #counts the NaN values to confirm the change

Number of zeros in 'Age' before replacement: 0
Number of NaN values in 'Age' after replacement: 24

df['Salary'].fillna(df['Salary'].median(), inplace=True)
df['Salary']

0      5000.0
1      5000.0
2      2300.0
3      3000.0
4      4000.0
...
143    9024.0
144    8787.0
145    4034.0
146    5034.0
147    8202.0
Name: Salary, Length: 148, dtype: float64

df['Place'].fillna('Unknown', inplace=True)

```


```
df.isnull().sum()
```

```
Company    0  
Age        24  
Salary     0  
Place      0  
Country    0  
Gender     0  
dtype: int64
```

## Handle NaN Values

```
sns.distplot(df['Age'])
```

```
<Axes: xlabel='Age', ylabel='Density'>
```



Note : Since the data is slightly skewed median will be better option

```
df['Age'].fillna(df['Age'].median(), inplace=True)  
print(df.isnull().sum())
```

```

Company    0
Age        0
Salary     0
Place      0
Country    0
Gender     0
dtype: int64

df.duplicated()

0      False
1      False
2      False
3      False
4      False
...
142     False
143     False
145     False
146     False
147     False
Length: 144, dtype: bool

df.duplicated().sum()

0

```

## Numerical columns and Categorical columns identified

```

numerical_columns = df.select_dtypes(include=['number']).columns
categorical_columns = df.select_dtypes(include=['object']).columns

print('numerical_columns:', numerical_columns)

numerical_columns: Index(['Age', 'Salary', 'Gender'], dtype='object')

print('categorical_columns:', categorical_columns)

categorical_columns: Index(['Company', 'Place', 'Country'],
dtype='object')

df

```

	Company	Age	Salary	Place	Country	Gender
0	TCS	20.0	5000.0	Chennai	India	0
1	Infosys	30.0	5000.0	Mumbai	India	0
2	TCS	35.0	2300.0	Calcutta	India	0
3	Infosys	40.0	3000.0	Delhi	India	0
4	TCS	23.0	4000.0	Mumbai	India	0
...	...	...	...	...	...	...
143	TCS	33.0	9024.0	Calcutta	India	1

144	Infosys	22.0	8787.0	Calcutta	India	1
145	Infosys	44.0	4034.0	Delhi	India	1
146	TCS	33.0	5034.0	Mumbai	India	1
147	Infosys	22.0	8202.0	Cochin	India	0

[148 rows x 6 columns]

## Data Exploration: (Score : 1)

Explore the data, list down the unique values in each feature and find its length.

```
for column in df.columns:
    unique_values = df[column].unique()
    print(f"Unique values in {column}: {unique_values}")
    print(f"Number of unique values in {column}: {len(unique_values)}\n")
```

Unique values in Company\_Name: [4 2 0 6 5 1 3]

Number of unique values in Company\_Name: 7

Unique values in Age: [20.	30.	35.	40.
23.	30.48461538		
34.	45.	18.	22.
50.	21.	46.	36.
24.	25.	43.	19.
31.	44.	33.	17.
			54.
			]

Number of unique values in Age: 29

Unique values in Salary: [5000. 2300. 3000. 4000. 6000. 7000. 8000. 9000. 1089. 1234. 3030. 3045. 3184. 4824. 5835. 7084. 8943. 8345. 9284. 9876. 2034. 7654. 2934. 4034. 5034. 8202. 9024. 4345. 6544. 6543. 3234. 4324. 5435. 5555. 8787. 3454. 5654. 5009. 5098. 3033.]

Number of unique values in Salary: 40

Unique values in Place: ['Chennai' 'Mumbai' 'Calcutta' 'Delhi' 'Podicherry' 'Cochin' 'Unknown' 'Noida' 'Hyderabad' 'Bhopal' 'Nagpur' 'Pune']

Number of unique values in Place: 12

Unique values in Country: [0]

Number of unique values in Country: 1

Unique values in Gender: [0, 1]

Categories (2, int64): [0, 1]

Number of unique values in Gender: 2

Perform the statistical analysis and renaming of the columns.

```
print(df.describe()) # describe() - to get a summary of the statistics for numerical features
```

	Company_Name	Age	Salary	Country
count	144.000000	144.000000	144.000000	144.0
mean	2.500000	31.855823	5238.194444	0.0
std	1.797434	8.250046	2370.641804	0.0
min	0.000000	17.000000	1089.000000	0.0
25%	1.000000	23.750000	3045.000000	0.0
50%	2.000000	32.000000	5000.000000	0.0
75%	4.000000	36.000000	7084.000000	0.0
max	6.000000	54.000000	9876.000000	0.0

```
print(df.describe(include = 'object')) # For categorical features, use describe(include='object')
```

	Place
count	144
unique	12
top	Mumbai
freq	34

## Renaming Columns

```
df.rename(columns={
    'Company': 'Company_Name'
}, inplace=True)
```

```
df['Company_Name']
```

0	TCS
1	Infosys
2	TCS
3	Infosys
4	TCS
...	
143	TCS
144	Infosys
145	Infosys
146	TCS
147	Infosys

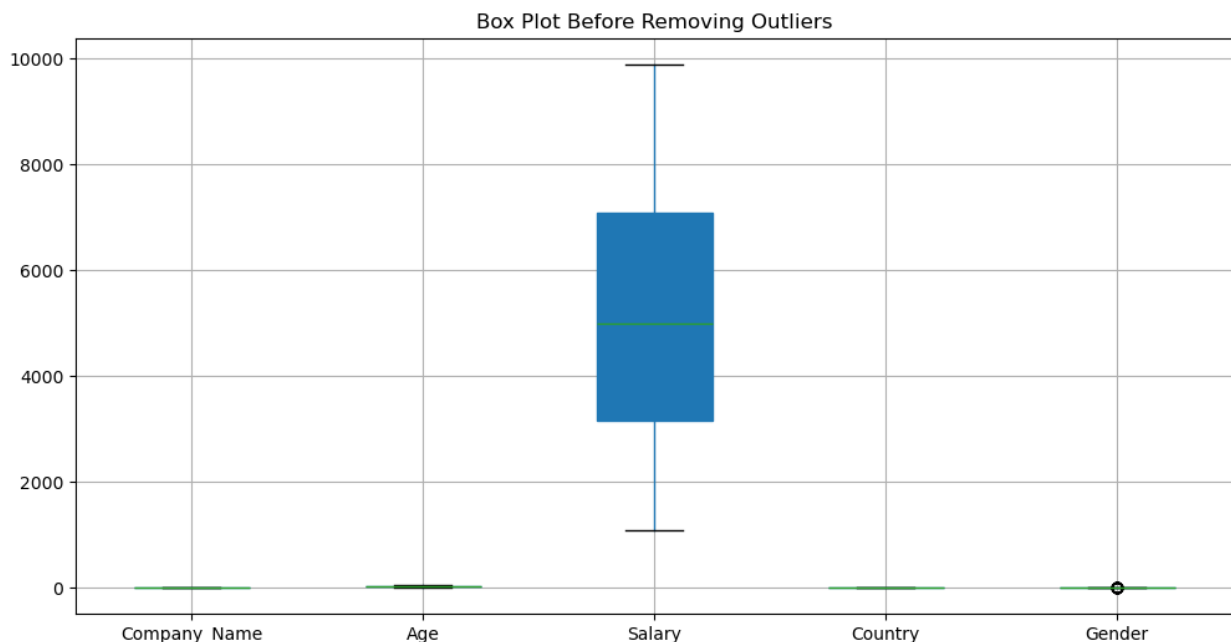
```
Name: Company_Name, Length: 148, dtype: object
```

## Convert Gender column to a categorical data type

```
df['Gender'] = df['Gender'].astype('category')
print(df['Gender'].dtype)
category
```

## Finding and Treating Outliers

```
# Box plot before removing outliers
df.select_dtypes(include='number').boxplot(figsize=(12, 6),
patch_artist=True)
plt.title("Box Plot Before Removing Outliers")
plt.show()
```



## Identify Outliers - IQR METHOD

```
for column in df.select_dtypes(include='number').columns:
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    outliers = df[(df[column] < lower_bound) | (df[column] >
upper_bound)]
    print(f"Number of outliers in {column}: {outliers.shape[0]}")
```



```

Number of outliers in Company_Name: 0
Number of outliers in Age: 0
Number of outliers in Salary: 0
Number of outliers in Country: 0
Number of outliers in Gender: 33

```


## Removed\_Capping

```

# Capping outliers to the lower and upper bounds
df[column] = df[column].clip(lower=lower_bound, upper=upper_bound)

# Box plot after capping outliers
df.select_dtypes(include='number').boxplot(figsize=(12, 6),
patch_artist=True)
plt.title("Box Plot After Capping Outliers")
plt.show()

```



## Data Analysis: (Score : 2)

Filter the data with age >40 and salary<5000

```

filtered_data = df[(df['Age'] > 40) & (df['Salary'] < 5000)]
filtered_data

```


	Company_Name	Age	Salary	Place	Country	Gender
21	2	50.0	3184.0	Delhi	0	0
32	2	45.0	4034.0	Calcutta	0	0
39	2	41.0	3000.0	Mumbai	0	0
50	2	41.0	3000.0	Chennai	0	0

57	2	51.0	3184.0	Hyderabad	0	0
68	2	43.0	4034.0	Mumbai	0	0
75	2	44.0	3000.0	Cochin	0	0
86	2	41.0	3000.0	Delhi	0	0
93	2	54.0	3184.0	Mumbai	0	0
104	2	44.0	4034.0	Delhi	0	0
122	2	44.0	3234.0	Mumbai	0	0
129	2	50.0	3184.0	Calcutta	0	0
138	0	44.0	3033.0	Cochin	0	0
140	2	44.0	4034.0	Hyderabad	0	0
145	2	44.0	4034.0	Delhi	0	1

## Plot the chart with age and salary

```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Salary', data=filtered_data)
plt.title("Age vs. Salary (Age > 40 and Salary < 5000)")
plt.xlabel("Age")
plt.ylabel("Salary")
plt.show()
```




Count the number of people from each place and represent it visually

```
# Count the number of people from each place
place_counts = df['Place'].value_counts()
place_counts

Place
Mumbai      34
Calcutta     32
Chennai      14
Delhi        14
Unknown      14
Cochin       13
Noida         8
Hyderabad    8
Podicherry   3
Pune          2
Bhopal        1
Nagpur         1
Name: count, dtype: int64

# Plot
plt.figure(figsize=(10, 6))
sns.barplot(x=place_counts.index, y=place_counts.values)
plt.title("Number of People from Each Place")
plt.xlabel("Place")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.show()
```



Data Encoding: (Score : 2)

Convert categorical variables into numerical representations using techniques such as one-hot encoding, label encoding, making them suitable for analysis by machine learning algorithms.

```
numerical_columns = df.select_dtypes(include=['number']).columns  
categorical_columns = df.select_dtypes(include=['object']).columns
```

Numerical Columns

```
print('numerical_columns:', numerical_columns)  
numerical_columns: Index(['Age', 'Salary', 'Gender'], dtype='object')
```

Categorical Columns

```
print('categorical_columns:', categorical_columns)
```

```
categorical_columns: Index(['Company_Name', 'Place', 'Country'],
dtype='object')
```

## Label Encoding

```
Label_Encoder=LabelEncoder() #label encoding convert categorical to numerical
```

```
df['Company_Name'] = Label_Encoder.fit_transform(df['Company_Name'])
df['Country'] = Label_Encoder.fit_transform(df['Country'])
```

```
print(df[['Company_Name']].head())
```

	Company_Name
0	4
1	2
2	4
3	2
4	4

```
print(df[['Country']].head())
```

	Country
0	0
1	0
2	0
3	0
4	0

## One-Hot Encoding

```
from sklearn.preprocessing import OneHotEncoder
```

```
# Assuming 'df' is your original DataFrame and you want to encode the 'Place' column
```

```
oneHot = OneHotEncoder(sparse_output=False) # Use sparse=False to return a dense array
```

```
# Apply OneHotEncoder to the 'Place' column
```

```
place_encoded = oneHot.fit_transform(df[['Place']])
```

```
# Get the column names for the one-hot encoded columns
```

```
place_columns = oneHot.get_feature_names_out(['Place'])
```

```
# Drop the original 'Place' column and concatenate the one-hot encoded columns
```

```
df_onehot = pd.concat([
    df.drop('Place', axis=1), # Drop the original 'Place' column
    pd.DataFrame(place_encoded, columns=place_columns) # Add the one-hot encoded columns
])
```

```
], axis=1)
```

```
df_onehot.head()
```

	Company_Name	Age	Salary	Country	Gender	Place_Bhopal	
0	Place_Calcutta \	4	20.0	5000.0	0	0	0.0
1		2	30.0	5000.0	0	0	0.0
2		4	35.0	2300.0	0	0	0.0
3		2	40.0	3000.0	0	0	0.0
4		4	23.0	4000.0	0	0	0.0

	Place_Chennai	Place_Cochin	Place_Delhi	Place_Hyderabad	
0	Place_Mumbai \	1.0	0.0	0.0	0.0
1		0.0	0.0	0.0	0.0
2		0.0	0.0	0.0	0.0
3		0.0	0.0	1.0	0.0
4		0.0	0.0	0.0	0.0

	Place_Nagpur	Place_Noida	Place_Podicherry	Place_Pune	
0	Place_Unknown	0.0	0.0	0.0	0.0
1		0.0	0.0	0.0	0.0
2		0.0	0.0	0.0	0.0
3		0.0	0.0	0.0	0.0
4		0.0	0.0	0.0	0.0

y is target variable (dependent variable)

```
y=df_onehot['Salary']
```

```
y
```

```
0      5000.0
1      5000.0
2      2300.0
3      3000.0
4      4000.0
```

```
...
143     9024.0
144     8787.0
145     4034.0
146     5034.0
147     8202.0
```

```
Name: Salary, Length: 148, dtype: float64
```

```
# x = df_onehot.drop(['Salary', 'Company_Name', 'Country'], axis=1)
```

```
x = df_onehot.drop(['Salary'], axis=1)
```

```
x
```

	Company_Name	Age	Country	Gender	Place_Bhopal	Place_Calcutta
0	4	20.0	0	0	0.0	0.0
1	2	30.0	0	0	0.0	0.0
2	4	35.0	0	0	0.0	1.0
3	2	40.0	0	0	0.0	0.0
4	4	23.0	0	0	0.0	0.0
..	...	...	...	...	...	...
143	4	33.0	0	1	0.0	1.0
144	2	22.0	0	1	0.0	1.0
145	2	44.0	0	1	0.0	0.0
146	4	33.0	0	1	0.0	0.0
147	2	22.0	0	0	0.0	0.0

	Place_Chennai	Place_Cochin	Place_Delhi	Place_Hyderabad
0	1.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0

0.0				
3	0.0	0.0	1.0	0.0
0.0				
4	0.0	0.0	0.0	0.0
1.0				
..	...	...	...	...
...				
143	0.0	0.0	0.0	0.0
0.0				
144	0.0	0.0	0.0	0.0
0.0				
145	0.0	0.0	1.0	0.0
0.0				
146	0.0	0.0	0.0	0.0
1.0				
147	0.0	1.0	0.0	0.0
0.0				

	Place_Nagpur	Place_Noida	Place_Podicherry	Place_Pune
Place_Unknown				
0	0.0	0.0	0.0	0.0
0.0				
1	0.0	0.0	0.0	0.0
0.0				
2	0.0	0.0	0.0	0.0
0.0				
3	0.0	0.0	0.0	0.0
0.0				
4	0.0	0.0	0.0	0.0
0.0				
..	...	...	...	...
...				
143	0.0	0.0	0.0	0.0
0.0				
144	0.0	0.0	0.0	0.0
0.0				
145	0.0	0.0	0.0	0.0
0.0				
146	0.0	0.0	0.0	0.0
0.0				
147	0.0	0.0	0.0	0.0
0.0				

[148 rows x 16 columns]

```
print(x.isnull().sum()) # Should show 0 if no `NaN` values exist
print(y.isnull().sum()) # Should also show 0 if no `NaN` values exist
```

Company_Name	0
Age	0



```

Country          0
Gender           0
Place_Bhopal     0
Place_Calcutta   0
Place_Chennai    0
Place_Cochin     0
Place_Delhi      0
Place_Hyderabad  0
Place_Mumbai     0
Place_Nagpur     0
Place_Noida      0
Place_Podicherry 0
Place_Pune       0
Place_Unknown    0
dtype: int64
0

```

## Feature Scaling: (Score : 2)

After the process of encoding, perform the scaling of the features using standardscaler and minmaxscaler.

## Feature Selection Using VarianceThreshold

```

# 1.1 Variance Threshold
var_threshold = VarianceThreshold(threshold=0.1)
X_var = var_threshold.fit_transform(x)
var_selected = x.columns[var_threshold.get_support()].tolist()

print("1. Filter Methods Results:")
print("\na) Variance Threshold")
print(f"Features selected: {len(var_selected)}")
print("Selected features:", var_selected[:5], "...")

```

1. Filter Methods Results:

a) Variance Threshold

Features selected: 5

Selected features: ['Company\_Name', 'Age', 'Gender', 'Place\_Calcutta', 'Place\_Mumbai'] ...

## scaling

```

# Create scalers
standard_scaler = StandardScaler() #standardized : entire data into
standard form
minmax_scaler = MinMaxScaler()     # minmax : entire data into
normalized form

```

```

# Apply different scaling methods
X_standardized = standard_scaler.fit_transform(X_var)
X_normalized = minmax_scaler.fit_transform(X_var)

# Convert the scaled arrays back into DataFrames with the original
column names
X_standardized_df = pd.DataFrame(X_standardized, columns=var_selected)
X_normalized_df = pd.DataFrame(X_normalized, columns=var_selected)

# now we can visualize or use these DataFrames
print(X_standardized_df.head())
print(X_normalized_df.head())

```

	Company_Name	Age	Gender	Place_Calcutta	Place_Mumbai
0	0.848436	-1.471033	-0.535683	-0.535683	-0.577350
1	-0.272712	-0.258148	-0.535683	-0.535683	1.732051
2	0.848436	0.348295	-0.535683	1.866775	-0.577350
3	-0.272712	0.954737	-0.535683	-0.535683	-0.577350
4	0.848436	-1.107168	-0.535683	-0.535683	1.732051

	Company_Name	Age	Gender	Place_Calcutta	Place_Mumbai
0	0.666667	0.081081	0.0	0.0	0.0
1	0.333333	0.351351	0.0	0.0	1.0
2	0.666667	0.486486	0.0	1.0	0.0
3	0.333333	0.621622	0.0	0.0	0.0
4	0.666667	0.162162	0.0	0.0	1.0

```

correlation = X_standardized_df.corr()
sns.heatmap(correlation, annot=True, cmap='coolwarm')
plt.title("Correlation Matrix of Standardized Features")
plt.show()

```

